

Advanced Time Series Analysis

State Space Models

Anshul Thakur & Prof. David Clifton

Computational Health Informatics Lab, IBME
Department of Engineering Science
University of Oxford

Table of Contents



State Space Models

Selective State Space Models

Further Reading

Table of Contents



State Space Models

Selective State Space Models

Further Reading

Structured State Space Model: A Linear RNN?

- SSMs use **linear state transitions** rather than non-linear hidden states
- **Time- and input-invariance:** rules governing state transitions and output generation do not change over the sequence

Discrete Linear SSM

At time step t the system follows:

$$\mathbf{h}_t = \mathbf{A} \mathbf{h}_{t-1} + \mathbf{B} \mathbf{x}_t \quad (\text{state update})$$

$$\mathbf{y}_t = \mathbf{C} \mathbf{h}_t + \mathbf{D} \mathbf{x}_t \quad (\text{output})$$

Notations:

$$\mathbf{x}_t \in \mathbb{R}^{d_{\text{in}}}, \mathbf{h}_t \in \mathbb{R}^n, \mathbf{y}_t \in \mathbb{R}^{d_{\text{out}}}, \mathbf{A} \in \mathbb{R}^{n \times n}, \mathbf{B} \in \mathbb{R}^{n \times d_{\text{in}}}, \mathbf{C} \in \mathbb{R}^{d_{\text{out}} \times n}, \mathbf{D} \in \mathbb{R}^{d_{\text{out}} \times d_{\text{in}}}.$$

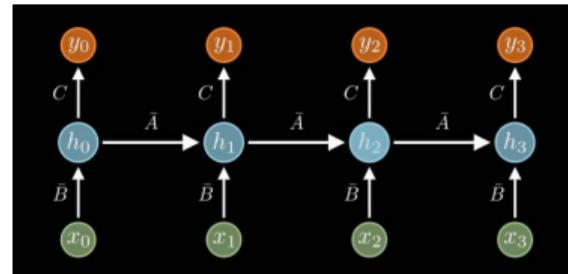


Figure: Discrete Linear SSM.

From Recurrence to Convolution

Discrete SSM:

$$\mathbf{h}_t = \mathbf{A}\mathbf{h}_{t-1} + \mathbf{B}\mathbf{x}_t, \quad \mathbf{y}_t = \mathbf{C}\mathbf{h}_t + \mathbf{D}\mathbf{x}_t$$

Unroll the recurrence:

$$\mathbf{h}_1 = \mathbf{B}\mathbf{x}_1,$$

$$\mathbf{h}_2 = \mathbf{A}\mathbf{h}_1 + \mathbf{B}\mathbf{x}_2 = \mathbf{A}\mathbf{B}\mathbf{x}_1 + \mathbf{B}\mathbf{x}_2,$$

$$\mathbf{h}_3 = \mathbf{A}\mathbf{h}_2 + \mathbf{B}\mathbf{x}_3 = \mathbf{A}^2\mathbf{B}\mathbf{x}_1 + \mathbf{A}\mathbf{B}\mathbf{x}_2 + \mathbf{B}\mathbf{x}_3,$$

$$\mathbf{h}_t = \sum_{k=0}^{t-1} \mathbf{A}^k \mathbf{B} \mathbf{x}_{t-k}.$$

Substitute into output:

$$\mathbf{y}_t = \mathbf{C}\mathbf{h}_t + \mathbf{D}\mathbf{x}_t = \sum_{k=0}^{t-1} \mathbf{C}\mathbf{A}^k \mathbf{B} \mathbf{x}_{t-k} + \mathbf{D}\mathbf{x}_t.$$

Convolution:



$$K[k] = \begin{cases} \mathbf{D} + \mathbf{C}\mathbf{B}, & k = 0, \\ \mathbf{C}\mathbf{A}^k \mathbf{B}, & k > 0, \end{cases}$$

$$\mathbf{y}_t = \sum_{k=0}^t K[k] \mathbf{x}_{t-k}$$

\iff

$$\mathbf{y} = K * \mathbf{x}$$

Intuition:

- ▶ Each \mathbf{A}^k describes how information propagates over k steps.
- ▶ $K[k]$ is the system's *impulse response* (weight for input lag k).
- ▶ Convolution sums all lagged contributions in parallel.

Practical Implications



- ▶ **A, B, C, and D** are learned end-to-end. The kernel weights $K[k]$ are generated from these matrices – not learned independently as in CNNs.
- ▶ **Pre-computation:** Build $K[0:T-1]$ once (using powers of **A**), then apply conv1d or FFT convolution: $\mathbf{y} = K * \mathbf{x}$. No Python-level time loop – full GPU parallelism.
- ▶ **Efficiency: RNNs:** $O(T)$ sequential steps. **SSMs:** $O(T \log T)$ parallel convolution – much faster on GPUs, enabling efficient modeling of very long sequences (audio, sensors, text).
- ▶ **Modeling benefit:** Combines **short-term** direct effects (via **D**) and **long-range** memory (via \mathbf{A}^k).

Table of Contents



State Space Models

Selective State Space Models

Further Reading

Structured vs Selective Models



Algorithm 1 SSM (S4)

Input: $x : (\mathbf{B}, \mathbf{L}, \mathbf{D})$

Output: $y : (\mathbf{B}, \mathbf{L}, \mathbf{D})$

1: $\mathbf{A} : (\mathbf{D}, \mathbb{N}) \leftarrow \text{Parameter}$

▷ Represents structured $N \times N$ matrix

2: $\mathbf{B} : (\mathbf{D}, \mathbb{N}) \leftarrow \text{Parameter}$

3: $\mathbf{C} : (\mathbf{D}, \mathbb{N}) \leftarrow \text{Parameter}$

4: $\Delta : (\mathbf{D}) \leftarrow \tau_\Delta(\text{Parameter})$

5: $\bar{\mathbf{A}}, \bar{\mathbf{B}} : (\mathbf{D}, \mathbb{N}) \leftarrow \text{discretize}(\Delta, \mathbf{A}, \mathbf{B})$

6: $y \leftarrow \text{SSM}(\bar{\mathbf{A}}, \bar{\mathbf{B}}, \mathbf{C})(x)$

▷ Time-invariant: recurrence or convolution

7: **return** y

$$\bar{\mathbf{A}} = e^{\Delta \mathbf{A}} \quad \bar{\mathbf{B}} = (e^{\Delta \mathbf{A}} - \mathbf{I}) \mathbf{A}^{-1} \mathbf{B}$$

Algorithm 2 SSM + Selection (S6)

Input: $x : (\mathbf{B}, \mathbf{L}, \mathbf{D})$

Output: $y : (\mathbf{B}, \mathbf{L}, \mathbf{D})$

1: $\mathbf{A} : (\mathbf{D}, \mathbb{N}) \leftarrow \text{Parameter}$

▷ Represents structured $N \times N$ matrix

2: $\mathbf{B} : (\mathbf{B}, \mathbf{L}, \mathbb{N}) \leftarrow s_B(x)$

3: $\mathbf{C} : (\mathbf{B}, \mathbf{L}, \mathbb{N}) \leftarrow s_C(x)$

4: $\Delta : (\mathbf{B}, \mathbf{L}, \mathbf{D}) \leftarrow \tau_\Delta(\text{Parameter} + s_\Delta(x))$

► 5: $\bar{\mathbf{A}}, \bar{\mathbf{B}} : (\mathbf{B}, \mathbf{L}, \mathbf{D}, \mathbb{N}) \leftarrow \text{discretize}(\Delta, \mathbf{A}, \mathbf{B})$

6: $y \leftarrow \text{SSM}(\bar{\mathbf{A}}, \bar{\mathbf{B}}, \mathbf{C})(x)$

▷ Time-varying: recurrence (*scan*) only

7: **return** y

$$s_B(\mathbf{x}) = \text{Linear}_N(\mathbf{x}) \quad s_\Delta(x) = \text{Broadcast}_D(\text{Linear}_1(x)).$$

$$\bar{\mathbf{A}}_{b,l,d,:,:} = e^{\Delta_{b,l,d} \mathbf{A}} \in \mathbb{R}^{N \times N}$$

$$\bar{\mathbf{B}}_{b,l,d,:} = (e^{\Delta_{b,l,d} \mathbf{A}} - \mathbf{I}) \mathbf{A}^{-1} \mathbf{B}_{b,l,:} \in \mathbb{R}^N$$

Scanning over Recurrence



Initial array

| | | | | | |
|---|---|---|----|---|---|
| 9 | 6 | 7 | 10 | 8 | 7 |
|---|---|---|----|---|---|

Prefix-Sum

| | | | | | |
|---|----|----|----|----|----|
| 9 | 15 | 22 | 32 | 40 | 47 |
|---|----|----|----|----|----|

$$h_t = \bar{A}h_{t-1} + \bar{B}x_t$$

The recurrent formula of the SSM model can also be thought of as a scan operation, in which each state is the sum of the previous state and the current input.

Model input

| | | | | | |
|-------|-------|-------|-------|-------|-------|
| x_0 | x_1 | x_2 | x_3 | x_4 | x_5 |
|-------|-------|-------|-------|-------|-------|

Scan output

| | | | | | |
|-------|-------|-------|-------|-------|-------|
| h_0 | h_1 | h_2 | h_3 | h_4 | h_5 |
|-------|-------|-------|-------|-------|-------|

- ▶ Scanning can be parallelized.
- ▶ Model selective space models such as Mamba use effective parallel scans.

Table of Contents



State Space Models

Selective State Space Models

Further Reading

Further Reading



- ▶ **Mamba & Mamba V2:** State Space Models enabling input-dependent dynamics
- ▶ **S4 / S5 Families:** Structured State Space Models introducing stable, parallelizable sequence representations.
- ▶ **Hybrid Models:** Combining SSMs with Transformers or Convolutional layers for multimodal and hierarchical tasks.